

Introduction to Computer Vision

Assignment #4

Due: Dec.-12 (Mon.) (before 11:59pm) No late submission allowed.

Instruction

- Submit your source codes, report, and descriptor in a single compressed file “CV_A4_StudentID.zip” to iCampus.
- Python 3.7 or higher / OpenCV 3.4 or higher will be used to execute your submitted codes.
- You can submit at most 7 files.
 - Mandatory files: *A4_StudentID.des*, *A4_compute_descriptors.py*, *A4_report.pdf*
- Any work that you turn in should be your own.

Part #1. Image Retrieval [100 pts]

The requirements will be evaluated based on your descriptor file, ‘*A4_StudentID.des*’

1-1. Dataset: a subset of UKBench

- Pre-extracted SIFT features of 1,000 images, *00000.sift* – *00999.sift*, are given by a compressed file, *CV_A4_Feats.zip*. They should be placed in the subdirectory ‘./feats/’ from your *.py* files. Your source codes should refer them by a relative directory address.
- For each image, there are 4 relevant items including itself. Those 4 images are ground truth retrieval results.
- Each SIFT feature corresponding to an image (*00xxx.sift*) has the following binary file format:

$$\begin{bmatrix} v_{1,1} & v_{1,2} & \dots & v_{1,128} \\ v_{2,1} & v_{2,2} & \dots & v_{2,128} \\ \vdots & \vdots & \ddots & \vdots \\ v_{n,1} & v_{n,2} & \dots & v_{n,128} \end{bmatrix}$$

The image has n SIFT features and each SIFT descriptor is 128-dimensional vector. Note that, each value $v_{i,j}$ is stored as ‘unsigned char’ format (1 byte representing a range from 0 to 255).

- Note that, 100 sample features (*s00000.sift* – *s00099.sift*) and their corresponding images (*s00000.jpg* – *s00099.jpg*) are also provided by *CV_A4_Sample.zip*. You can utilize them to verify and debug your implementations.

1-2. Task: computing image descriptors for the retrieval task

- Your task is to compute image descriptors that reflect similarities among images by L_1 or L_2 distances. You should store the descriptors as the following binary format file:

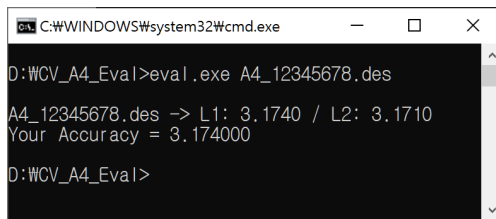
$\begin{matrix} N & D \\ d_{1,1} & d_{1,2} & \dots & d_{1,D} \\ d_{2,1} & d_{2,2} & \dots & d_{2,D} \\ \dots & & & \\ d_{N,1} & d_{N,2} & \dots & d_{N,D} \end{matrix}$

, where N and D are the number of images (=1000) and the dimensionality of each descriptor, respectively. And $d_{i,j}$ is j^{th} element of i^{th} descriptor. The required data types of variables are summarized as follows:

Variable	Type
N	Signed 32 bits integer (<i>int</i> in C/C++)
D	Signed 32 bits integer (<i>int</i> in C/C++)
$d_{i,j}$	32 bits floating point (<i>float</i> in C/C++)

Note that, the dimensionality of your descriptor cannot be greater than 1,024 (i.e. $D \leq 1024$). In other words, the size of your descriptor file '*A4_StudentID.des*' can be up to 4,096,008 bytes ($= 2 \times 4 + 1000 \times 1024$ bytes).

- b) You need to submit a code to produce the descriptors. Your script should run when they are located at any directory containing the '*./feats/*' sub-directory. In other words, you should specify the data path as a relative directory path in your implementation. Name your code as '*A4_compute_descriptors.py*'. If your algorithm utilizes codeword vectors based on the clustering (taking long computation time), you are asked to precompute/store and load them in your script. In this case, you should include these precomputed files in your submission package.
- c) Once you produced the descriptor file, you can evaluate your descriptors by using provided executable '*eval.exe*'. The program will evaluate your descriptors according to both L_1 and L_2 distances and take the higher one. It only works in Windows OS. You have to specify your descriptor file name as the following example:



```

C:\WINDOWS\system32\cmd.exe
D:\WCV_A4_Eval>eval.exe A4_12345678.des
A4_12345678.des -> L1: 3.1740 / L2: 3.1710
Your Accuracy = 3.174000
D:\WCV_A4_Eval>

```

Note that, the accuracy is ranging from 0 to 4.

- d) You are required to submit a report '*A4_report.pdf*' that describes your approach to compute the image descriptors. It should be longer than 1 page but shorter than 3 pages.
- e) It is strongly prohibited to produce image descriptors by abusing the characteristics of the dataset. One typical example is setting $d_1 = d_2 = d_3 = d_4 = [1,0,0, \dots, 0]$, $d_5 = d_6 = d_7 = d_8 = [0,1,0, \dots, 0]$.